

XGBoost

September 9, 2020

STAT 685: Dr. Suojin Wang's Group

Modeling Seoul Bike Sharing Demand - 003 XGBoosting Study considering Time

Bai Zou, Nam Tran

1 Key Takeaways

- Comparing with model assuming independent observation, using anchor date in model training and testing shows a lower performance.
- Information related to time of the year is missing in model training data when an anchor date is set.
- Adding time related attributes has limited improvement.
- Next step:
 - More parameter tuning (Baseline performance is lower than referred literature)
 - Study effect with different anchor day setting and prediction horizon. E.g., predict last few days hourly demand by end of the season; predict next few hours demand in a random day by knowing all information before; etc.

2 Data Process

- Load data and check data format
- Add dummy variables and convert qualitative variable to quantitative
- Randomly select 25% data as testing data

```
[1]: from numpy import loadtxt
import xgboost as xgb
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import os
import time
import pandas as pd
import matplotlib.pyplot as plt
```

```
[20]: ## load data
```

```

fn = os.path.abspath(os.path.join(os.path.dirname(os.getcwd()), './data/
↳SeoulBikeData.csv'))
colNames = ["Date", "RentedBikeCount", "Hour", "Temp", "Humidity", "WindSpeed", "
↳"Visibility",
            "DewPointTemp", "SolarRadiation", "Rainfall", "Snowfall", "Seasons", "
↳"Holiday", "FunctionalDay"]
dat = pd.read_csv(fn, encoding="ISO-8859-1")
dat.columns = colNames
dat = dat.astype(
    {"Date": str,
     "RentedBikeCount": int,
     "Hour": float,
     "Temp": float,
     "Humidity": float,
     "WindSpeed": float,
     "Visibility": float,
     "DewPointTemp": float,
     "SolarRadiation": float,
     "Rainfall": float,
     "Snowfall": float,
     "Seasons": str,
     "Holiday": str,
     "FunctionalDay": str})

# convert qualitative variable to quantitative
dat = pd.get_dummies(dat, columns=["Seasons"])
del dat['Seasons_Winter']
dat = pd.get_dummies(dat, columns=["Holiday"])
del dat['Holiday_Holiday']
dat = pd.get_dummies(dat, columns=["FunctionalDay"])
del dat['FunctionalDay_Yes']
dat.shape

```

[20]: (8760, 16)

[4]: dat.head(2)

```

[4]:
      Date  RentedBikeCount  Hour  Temp  Humidity  WindSpeed  Visibility \
0  01/12/2017             254   0.0  -5.2     37.0         2.2     2000.0
1  01/12/2017             204   1.0  -5.5     38.0         0.8     2000.0

      DewPointTemp  SolarRadiation  Rainfall  Snowfall  Seasons_Autumn \
0          -17.6             0.0         0.0         0.0             0
1          -17.6             0.0         0.0         0.0             0

      Seasons_Spring  Seasons_Summer  Holiday_No  Holiday  FunctionalDay_No
0                0                0            0          1                0

```

1 0 0 1 0

```
[5]: # split data into X and y
X = dat.iloc[:, 2:dat.shape[1]]
Y = dat.iloc[:,1]
print(f"X size {X.shape} and Y size {Y.shape}")
```

X size (8760, 14) and Y size (8760,)

3 Model Assuming Independent Observations

- Training and Testing data are selected randomly across the year.
- The 'Date' information is not used.

```
[6]: # split data into train and test sets
seed = 7
test_size = 0.25
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=test_size,
↳random_state=seed)
```

```
[7]: print(f"X_train size {X_train.shape}; y_train size {y_train.shape}")
```

X_train size (6570, 14); y_train size (6570,)

```
[8]: print(f"X_test size {X_test.shape}; y_test size {y_test.shape}")
```

X_test size (2190, 14); y_test size (2190,)

3.1 Simple XGBoost Model

- To start with, simply train the model with all training set from previous splitting.
- The training set prediction accuracy reaches 99.997 % but testing accuracy is at 85.487 %.
- The testing result is not improving after several iterations.
- The next step is to add CV and tune some parameters

```
[45]: %%capture
# fit model with training data
start_time = time.time()
model1 = xgb.XGBRegressor(n_estimators=1000)
model1.fit(X_train, y_train,
           eval_set=[(X_train, y_train), (X_test, y_test)],
           verbose=True)
```

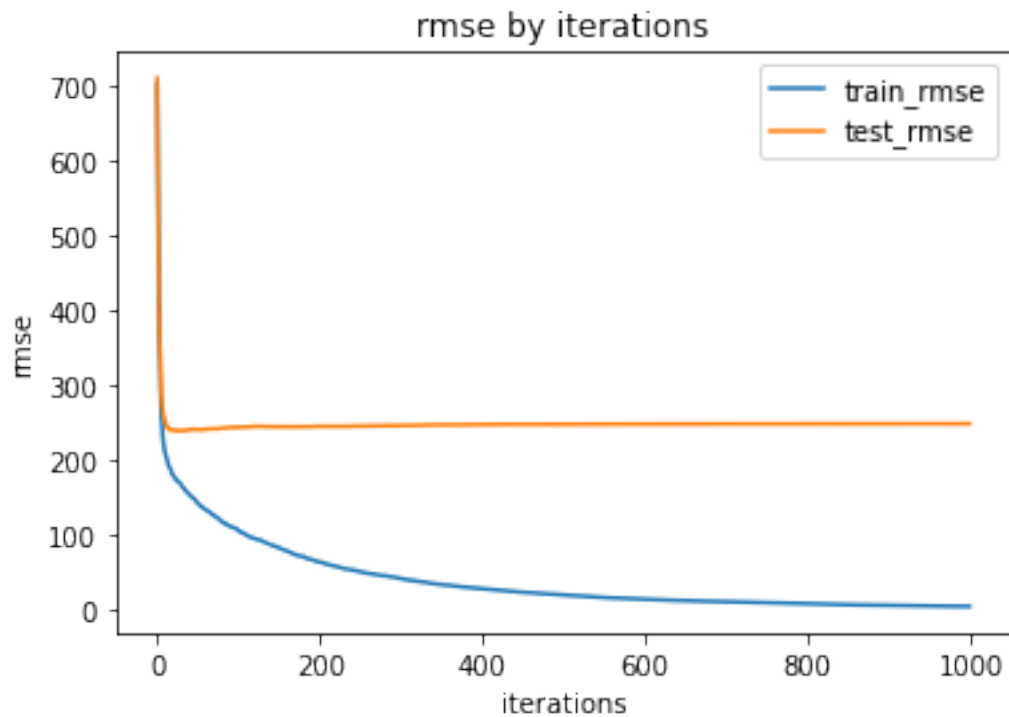
```
[46]: print(f"Model running for {(time.time()-start_time)/60:0.2f} min.")
```

Model running for 0.04 min.


```
[47]: evals_result1 = model1.eval_result()
```

```
[48]: def plot_validation(evals_result, metric='rmse', ax=None, title=None):  
    train_rmse = evals_result['validation_0'][metric]  
    test_rmse = evals_result['validation_1'][metric]  
    plot_dat = pd.DataFrame({'train_rmse': train_rmse,  
                             'test_rmse': test_rmse})  
    title = f"{metric} by iterations" if title is None else title  
    plot_dat.plot(xlabel="iterations", ylabel=metric, ax=ax, title=title)
```

```
[49]: plot_validation(evals_result1)
```



```
[232]: _ = xgb.plot_importance(model1, height=0.9)
```



```
[145]: score = model1.score(X_train, y_train)
print(f"Training score: {score*100:.3f} %")
```

Training score: 99.997 %

```
[146]: score = model1.score(X_test, y_test)
print(f"Testing score: {score*100:.3f} %")
```

Testing score: 85.487 %

3.2 Model with GridSearchCV

```
[35]: from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
```

```
[279]: %%capture
start_time = time.time()
parameters = {'nthread':[4], #when use hyperthread, xgboost may become slower
              'objective':['reg:squarederror'],
              'learning_rate': [.03, .07, 0.3], #so called `eta` value
              'max_depth': [3, 5, 7],
              'min_child_weight': [4],
              'silent': [1],
              'subsample': [0.7],
              'colsample_bytree': [0.7],
```

```

        'n_estimators': [500]}


xgb_grid = GridSearchCV(xgb.XGBRegressor(),
                        parameters,
                        cv = 10,
                        n_jobs = 5,
                        verbose=True)
model2 = xgb_grid.fit(X_train, y_train,
                     eval_set=[(X_train, y_train), (X_test, y_test)],
                     verbose=True)
model2 = model2.best_estimator_

```


```
[284]: print(f"Model running for {(time.time()-start_time)/60:0.2f} min.")
```

Model running for 2.48 min.

```
[285]: evals_result2 = model2.evals_result()
       plot_validation(evals_result2)
```



```
[286]: _ = xgb.plot_importance(model2, height=0.9)
```



```
[287]: score = model2.score(X_train, y_train)
print(f"Training score: {score*100:.3f} %")
```

Training score: 92.766 %

```
[288]: score = xgb_grid.cv_results_['mean_test_score'].mean()
print(f"CV mean test score: {score*100:.3f} %")
```

CV mean test score: 86.589 %

```
[289]: score = model2.score(X_test, y_test)
print(f"Testing score: {score*100:.3f} %")
```

Testing score: 87.148 %

4 Model Tested with Anchor Date

- In this time series data, an anchor date is defined to split training and testing data.
- The model is trained based on data before anchor date and used to predict observations after anchor date to evaluate accuracy.
- Time series attribution is not included in this model.
- The testing results show:
 - In general lower accuracy compared with random testing data selection
 - The training data lacks information for season Autumn.

- Non-function day attribute is not applied well based on prediction comparison plots. Non-function day record appears mainly during months 9, 10 and 11. Therefore, testing month 11 has the best prediction accuracy as month 9 and 10 information is included in training data.

4.1 Training & Testing Data Splitting with Anchor Date

- In this time series data, observations are based by hour and date.
- 25% testing data equals to around 91 days or 3 months.
- Split the data into 3 sets of training and testing data.

```
[21]: dat['Date'] = pd.to_datetime(dat['Date'], format="%d/%m/%Y")
```

```
[25]: dat['Date'].dt.month.unique()
```

```
[25]: array([12, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11])
```

```
[115]: dat_set = {}

for test_month in [9, 10, 11]:
    train_dat = dat[(dat['Date'].dt.month < test_month) | (dat['Date'].dt.year == 2017)]
    test_dat = dat[dat['Date'].dt.month == test_month]
    print(f"Training Data Size {train_dat.shape}; Testing Data Size {test_dat.shape}")
    sub_X_train = train_dat.iloc[:, 2:train_dat.shape[1]]
    sub_Y_train = train_dat.iloc[:, 1]
    sub_X_test = test_dat.iloc[:, 2:test_dat.shape[1]]
    sub_Y_test = test_dat.iloc[:, 1]
    dat_set[test_month] = (sub_X_train, sub_Y_train, sub_X_test, sub_Y_test)
```

Training Data Size (6576, 16); Testing Data Size (720, 16)

Training Data Size (7296, 16); Testing Data Size (744, 16)

Training Data Size (8040, 16); Testing Data Size (720, 16)

4.2 Model Training and Testing

```
[97]: def cv_xgb_train(X_training, Y_training, X_testing, Y_testing):
    param = {'nthread': [4], #when use hyperthread, xgboost may become slower
             'objective': ['reg:squarederror'],
             'learning_rate': [.03, .07, 0.3], #so called `eta` value
             'max_depth': [3, 5, 7],
             'min_child_weight': [4],
             'silent': [1],
             'subsample': [0.7],
             'colsample_bytree': [0.7],
```



```

        'n_estimators': [500]}

x_g = GridSearchCV(xgb.XGBRegressor(),
                  param,
                  cv = 10,
                  n_jobs = 5,
                  verbose=True)
mod = x_g.fit(X_training,
             Y_training,
             eval_set=[(X_training, Y_training), (X_testing, Y_testing)],
             verbose=True)
mod = mod.best_estimator_
return mod, x_g

```

```

[98]: %%capture
start_time = time.time()
res_summary = {}
res_mod = {}
for k, v in dat_set.items():
    sub_X_train, sub_Y_train, sub_X_test, sub_Y_test = v
    fit_mod, trained_grid = cv_xgb_train(sub_X_train, sub_Y_train, sub_X_test,
    ↪sub_Y_test)
    # track mod
    res_mod[k] = fit_mod
    # track metric
    res_summary[k] = {}
    res_summary[k]['train_score'] = fit_mod.score(sub_X_train, sub_Y_train)
    res_summary[k]['mean_cv_score'] = trained_grid.
    ↪cv_results_['mean_test_score'].mean()
    res_summary[k]['test_score'] = fit_mod.score(sub_X_test, sub_Y_test)

```

```

[100]: print(f"Model running for {(time.time()-start_time)/60:0.2f} min.")

```

Model running for 7.92 min.

```

[99]: pd.DataFrame.from_dict(res_summary)

```

```


[99]:
          9          10          11
train_score    0.933605  0.928365  0.924187
mean_cv_score  0.624540  0.628203  0.622360
test_score     0.775022  0.606972  0.763277

```


```

[101]: fig, ax = plt.subplots(nrows=1, ncols=3, figsize=(21,5))
for k, mod in res_mod.items():
    evals_res = mod.evals_result()
    plot_validation(evals_res, metric='rmse', ax=ax[k-9], title=f"Testing Month_
    ↪{k}")

```




```
[102]: fig, ax = plt.subplots(nrows=1, ncols=3, figsize=(21,5))
for k, mod in res_mod.items():
    xgb.plot_importance(mod, height=0.9, ax=ax[k-9], title=f"Testing Month {k}")
```



4.3 Prediction Accuracy by Time

```
[107]: def plot_prediction(mod, sub_X_test, sub_Y_test, ax=None, title=None):
    y_pred = mod.predict(sub_X_test)
    plot_dat = pd.DataFrame({"Predicted_Y_Teset": y_pred,
                             "Recorded_Y_Teset": sub_Y_test})
    plot_dat.plot(xlabel="time", ylabel="count", ax=ax, title=title)
```

```
[118]: fig, ax = plt.subplots(nrows=3, ncols=1, figsize=(21,15))
for k, mod in res_mod.items():
    sub_X_train, sub_Y_train, sub_X_test, sub_Y_test = dat_set[k]
    plot_prediction(mod, sub_X_test, sub_Y_test, ax=ax[k-9], title=f"Testing_
    ↪Month {k}")
```



4.4 Non-function Days

The plots above show a bad prediction on non-function days in month 9. By checking the non-function days across overall data: * 2 days before September 2018 * 4 days in September 2018 * 4 days in October 2018 * 3 days in November 2018

```
[138]: dat.groupby('FunctionalDay_No').agg({'Date': 'nunique'})
```

```
[138]:
```

	Date
FunctionalDay_No	
0	353
1	13

```
[139]: dat[dat['FunctionalDay_No']==1]['Date'].unique()
```

```
[139]: array(['2018-04-11T00:00:00.000000000', '2018-05-10T00:00:00.000000000',
          '2018-09-18T00:00:00.000000000', '2018-09-19T00:00:00.000000000',
          '2018-09-28T00:00:00.000000000', '2018-09-30T00:00:00.000000000',
          '2018-10-02T00:00:00.000000000', '2018-10-04T00:00:00.000000000',
          '2018-10-06T00:00:00.000000000', '2018-10-09T00:00:00.000000000',
          '2018-11-03T00:00:00.000000000', '2018-11-06T00:00:00.000000000',
```

```
'2018-11-09T00:00:00.000000000'], dtype='datetime64[ns]')
```

4.5 Model Detail

```
[155]: for k, mod in res_mod.items():  
        print(f"Testing month {k}: \n{mod}")
```

Testing month 9:

```
XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
             colsample_bynode=1, colsample_bytree=0.7, gamma=0, gpu_id=-1,  
             importance_type='gain', interaction_constraints='',  
             learning_rate=0.03, max_delta_step=0, max_depth=5,  
             min_child_weight=4, missing=nan, monotone_constraints=(),  
             n_estimators=500, n_jobs=4, nthread=4, num_parallel_tree=1,  
             random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,  
             silent=1, subsample=0.7, tree_method='exact',  
             validate_parameters=1, verbosity=None)
```

Testing month 10:

```
XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
             colsample_bynode=1, colsample_bytree=0.7, gamma=0, gpu_id=-1,  
             importance_type='gain', interaction_constraints='',  
             learning_rate=0.03, max_delta_step=0, max_depth=5,  
             min_child_weight=4, missing=nan, monotone_constraints=(),  
             n_estimators=500, n_jobs=4, nthread=4, num_parallel_tree=1,  
             random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,  
             silent=1, subsample=0.7, tree_method='exact',  
             validate_parameters=1, verbosity=None)
```

Testing month 11:

```
XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
             colsample_bynode=1, colsample_bytree=0.7, gamma=0, gpu_id=-1,  
             importance_type='gain', interaction_constraints='',  
             learning_rate=0.03, max_delta_step=0, max_depth=5,  
             min_child_weight=4, missing=nan, monotone_constraints=(),  
             n_estimators=500, n_jobs=4, nthread=4, num_parallel_tree=1,  
             random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,  
             silent=1, subsample=0.7, tree_method='exact',  
             validate_parameters=1, verbosity=None)
```

5 Model with Time Attributes

- There's some improvements in testing but along with slightly overfitting in training data.
- The added attributes ranked in the mid range in importance plots. Temp and Hour are still the two most important attributes.

```
[121]: dat2 = dat.copy()
dat2['Year'] = dat2['Date'].dt.year
dat2['Month'] = dat2['Date'].dt.month
dat2['Week'] = dat2['Date'].dt.week
dat2['Day'] = dat2['Date'].dt.day
dat2['DayOfWeek'] = dat2['Date'].dt.dayofweek
```

<ipython-input-121-1012b0438a4f>:4: FutureWarning: Series.dt.weekofyear and Series.dt.week have been deprecated. Please use Series.dt.isocalendar().week instead.

```
dat2['Week'] = dat2['Date'].dt.week
```

```
[122]: dat2.head(2)
```

```
[122]:      Date  RentedBikeCount  Hour  Temp  Humidity  WindSpeed  Visibility \
0 2017-12-01             254    0.0  -5.2     37.0         2.2     2000.0
1 2017-12-01             204    1.0  -5.5     38.0         0.8     2000.0

      DewPointTemp  SolarRadiation  Rainfall  ...  Seasons_Autumn  \
0             -17.6             0.0        0.0  ...              0
1             -17.6             0.0        0.0  ...              0

      Seasons_Spring  Seasons_Summer  Holiday_No  Holiday  FunctionalDay_No  Year  \
0                0                0            1         1                0  2017
1                0                0            1         1                0  2017

      Month  Week  Day  DayOfWeek
0        12    48    1          4
1        12    48    1          4
```

[2 rows x 21 columns]

```
[123]: dat_set2 = {}

for test_month in [9, 10, 11]:
    train_dat = dat2[(dat2['Date'].dt.month < test_month)|(dat2['Date'].dt.year_
    ↪== 2017)]
    test_dat = dat2[dat2['Date'].dt.month == test_month]
    print(f"Training Data Size {train_dat.shape}; Testing Data Size {test_dat.
    ↪shape}")
    sub_X_train = train_dat.iloc[:, 2:train_dat.shape[1]]
    sub_Y_train = train_dat.iloc[:, 1]
    sub_X_test = test_dat.iloc[:, 2:test_dat.shape[1]]
    sub_Y_test = test_dat.iloc[:, 1]
    dat_set2[test_month] = (sub_X_train, sub_Y_train, sub_X_test, sub_Y_test)
```

Training Data Size (6576, 21); Testing Data Size (720, 21)

Training Data Size (7296, 21); Testing Data Size (744, 21)
Training Data Size (8040, 21); Testing Data Size (720, 21)

```
[156]: def cv_xgb_train2(X_training, Y_training, X_testing, Y_testing):
    param = {'nthread':[4], #when use hyperthread, xgboost may become slower
             'objective':['reg:squarederror'],
             'learning_rate': [.03, .07, 0.3], #so called `eta` value
             'max_depth': [3, 5], # limiting max_depth to avoid_
             →overfitting
             'min_child_weight': [4],
             'silent': [1],
             'subsample': [0.7],
             'colsample_bytree': [0.7],
             'n_estimators': [500]}

    x_g = GridSearchCV(xgb.XGBRegressor(),
                       param,
                       cv = 10,
                       n_jobs = 5,
                       verbose=True)
    mod = x_g.fit(X_training,
                  Y_training,
                  eval_set=[(X_training, Y_training), (X_testing, Y_testing)],
                  verbose=True)
    mod = mod.best_estimator_
    return mod, x_g

[157]: %%capture
start_time = time.time()
res_summary2 = {}
res_mod2 = {}
for k, v in dat_set2.items():
    sub_X_train, sub_Y_train, sub_X_test, sub_Y_test = v
    fit_mod, trained_grid = cv_xgb_train2(sub_X_train, sub_Y_train, sub_X_test,
    →sub_Y_test)
    # track mod
    res_mod2[k] = fit_mod
    # track metric
    res_summary2[k] = {}
    res_summary2[k]['train_score'] = fit_mod.score(sub_X_train, sub_Y_train)
    res_summary2[k]['mean_cv_score'] = trained_grid.
    →cv_results_['mean_test_score'].mean()
    res_summary2[k]['test_score'] = fit_mod.score(sub_X_test, sub_Y_test)

[158]: print(f"Model running for {(time.time()-start_time)/60:0.2f} min.")
```


Model running for 9.74 min.

```
[159]: pd.DataFrame.from_dict(res_summary2) # TODO check negative values
```


```
[159]:
```

	9	10	11
train_score	0.968569	0.964197	0.981632
mean_cv_score	0.687866	0.575854	-0.235344
test_score	0.799063	0.714356	0.825770


```
[161]: fig, ax = plt.subplots(nrows=1, ncols=3, figsize=(21,5))
for k, mod in res_mod2.items():
    evals_res = mod.evals_result()
    plot_validation(evals_res, metric='rmse', ax=ax[k-9], title=f"Testing Month_{k}")
```



```
[162]: fig, ax = plt.subplots(nrows=1, ncols=3, figsize=(21,5))
for k, mod in res_mod2.items():
    xgb.plot_importance(mod, height=0.9, ax=ax[k-9], title=f"Testing Month {k}")
```



```
[163]: fig, ax = plt.subplots(nrows=3, ncols=1, figsize=(21,15))
for k, mod in res_mod2.items():
    sub_X_train, sub_Y_train, sub_X_test, sub_Y_test = dat_set2[k]
    plot_prediction(mod, sub_X_test, sub_Y_test, ax=ax[k-9], title=f"Testing_Month_{k}")
```



```
[160]: for k, mod in res_mod2.items():
        print(f"Testing month {k}: \n{mod}")
```

Testing month 9:

```
XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=0.7, gamma=0, gpu_id=-1,
              importance_type='gain', interaction_constraints='',
              learning_rate=0.03, max_delta_step=0, max_depth=5,
              min_child_weight=4, missing=nan, monotone_constraints='()',
              n_estimators=500, n_jobs=4, nthread=4, num_parallel_tree=1,
              random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
              silent=1, subsample=0.7, tree_method='exact',
              validate_parameters=1, verbosity=None)
```

Testing month 10:

```
XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=0.7, gamma=0, gpu_id=-1,
              importance_type='gain', interaction_constraints='',
              learning_rate=0.03, max_delta_step=0, max_depth=5,
              min_child_weight=4, missing=nan, monotone_constraints='()',
              n_estimators=500, n_jobs=4, nthread=4, num_parallel_tree=1,
              random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
              silent=1, subsample=0.7, tree_method='exact',
```



```
        validate_parameters=1, verbosity=None)
Testing month 11:
XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
             colsample_bynode=1, colsample_bytree=0.7, gamma=0, gpu_id=-1,
             importance_type='gain', interaction_constraints='',
             learning_rate=0.07, max_delta_step=0, max_depth=5,
             min_child_weight=4, missing=nan, monotone_constraints='()',
             n_estimators=500, n_jobs=4, nthread=4, num_parallel_tree=1,
             random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
             silent=1, subsample=0.7, tree_method='exact',
             validate_parameters=1, verbosity=None)
```